

# MINGYONG MA

3869 Miramar St, San Diego, US

☎ 858-539-6919 ✉ [m7ma@ucsd.edu](mailto:m7ma@ucsd.edu) [My Website](#) [github.com/mamioma](https://github.com/mamioma)

## Education

University of California, San Diego

*Master of Science in Computer Science*

Sep. 2022 – Dec 2023

*San Diego, California*

## Experience

Adobe

June 2023 – Sep 2022

*Software Engineer Intern*

*San Jose, California*

- Integrated **LLM** model **fine-tuning** and **inference** features to Adobe primary AI platform **Firefall**.
- Created a service that user can submit a **fine-tuning** task, and this request will be forwarded to **MMS** (Model Management System), which will download our code from remote artifactory **JFrog**, integrate into **MMS** and build a **Docker** container. After fine-tuning task is finished, it will save the fine-tuned model in **Azure Blob Storage**, and returns the results back to the user.
- Improved the fine-tuning API call from **blocking** to **asynchronous**. Users no longer need to wait for the fine-tuning results. They receive a **task\_id** instantly when submitting a fine-tuning task, which can be used to query **Firefall** for the results. This modification has reduced **overhead** and boosted **latency** by 90%.
- **Reduced the network I/O** from **13GB** to **32MB** per **inference** call. By utilizing **PEFT**, the based model is consistent for every fine-tuning job, thus is stored in **in-memory-buffer** of the **Docker** container, with only the **Lora layer** being stored in **Azure Blob Storage**. Therefore, only the Lora layer (32MB) instead of the entire model (13GB) is downloaded into Docker container.
- Used **Jmeter** for **load-testing**, able to generate **1600 TPS** (token per second) with **multi-threading**.
- Innovatively proposed how to fine-tune **LLaMa2-7b** on a **CPU**, which offers alternative choice to save cost. No need to run GPU entirely a day.
- Implemented using **REST API** that able to **CRUD** a task, and save it in **postgres DB** with **almebic** version control.

Amazon

June 2022 – August 2022

*Software engineer Intern*

*Shenzhen*

- Developed an **image processing** algorithm that combines **deep learning** techniques with the **Unsharp** algorithm, achieving superior results compared to the camera algorithm used in tablets. And evaluated the performance of the system using **MTF-50**.
- Demonstrated ability to compare the performance of algorithms by controlling the imaging device with **adb** and generating identical images with different image sharpening settings in the Amazon lab.
- Conducted an evaluation of our proposed algorithm using **Imatest** software in the Amazon lab, observing an increase in **MTF-50**, which showcases an improvement in image sharpness.

Lenovo

Nov 2021 – Feb 2022

*Data Analytic Intern*

*Beijing*

- Conducted **time series forecasting** to predict future sales of Lenovo's notebook products and tablets, utilizing Lenovo's historical sales data as well as data from other companies such as IDC and GFK.
- Increased the forecasting accuracy of the model by 4.2% by implementing machine learning algorithms such as **Prophet** and deep learning models like **LSTM** or **LSTM**.

## Projects

Database acceleration | *C++14, mutex*

March 2023

- Realized a Database index method utilizing **B+ Tree**, which shows **10 times faster** performance compared with **Hash index** or **file scan** on **range search**.
- Built a **Buffer Pool** on top of I/O layer, and realize **Buffer Replacement Policy** and **LRU clock algorithm**.
- Built a **B+ Tree** on top of Buffer Pool, supporting **CRUD** operation. Besides, it can save more than **50GB** data.

Distributed Cloud File System | *Go, gRPC*

January 2023

- Created a **fault tolerant cloud file storage** service called SurfStore (client and server communicating using **gRPC**).
- Stored and manage the block in different BlockStore using **Consistent Hashing Ring**.
- Ensured the MetaStore is fault-tolerant and stays consistent regardless of minority of server failures by **RAFT** protocol.

Operating System Implemetation | *Java*

Sep 2022

- Implemented **life cycle** of the OS process, **virtual memory** and **file system**.
- Created the **pageTable** for each user process, which maps the process's **virtual addresses** to **physical addresses**.
- Implement **demand paging**, **page replacement** to free up a physic page to handle **page faults**.